



WaxFrame

MANY MINDS, ONE REFINED RESULT.

Getting Started

First-run guide for new users

What is WaxFrame?



WaxFrame puts a team of AI Worker Bees on your document at the same time.

ChatGPT, Claude, Gemini, Grok, Mistral, Perplexity, Together AI — whichever AIs you have API keys for — review and refine your document in iterative rounds, while a designated **Builder** AI synthesizes the hive's feedback into the next version.

Where a single AI gives you one perspective, a hive surfaces disagreements, catches blind spots, and converges on a stronger result.

OVERVIEW

WaxFrame runs in your browser. **No backend, no install, no signup.** Your API keys, prompts, and documents stay in your browser's local storage. The whole app is one HTML page plus its assets — air-gap safe and AGPL-3.0 open source.

How a round works

- 1 You write or upload a starting document**
— or start from a template (23 ship in the box).
- 2 The hive reviews it in parallel**
— each Worker Bee returns a numbered list of suggestions.
- 3 The Builder synthesizes**
— it reads every reviewer's feedback, flags genuine disagreements as Conflicts for you to resolve, and produces the next refined version of the document.
- 4 You read the result, resolve any conflicts, and run another round**
— repeat until the hive converges on a result it can no longer materially improve.

Your first run in 5 steps

Total setup time on a fresh install: about 10 minutes. Most of that is signing up for one or two API keys.

1 Open WaxFrame in a desktop browser.

Open `index.html` (or visit waxframe.com) on a desktop or laptop. Minimum viewport is **1366 × 768 px**; 1600+ wide is recommended for the multi-panel work screen. Phones and tablets won't fit the workflow — you'll see a mobile landing page that points you to mobile-readable docs.

2 Add API keys for at least 2 AIs.

The first screen is **Your Worker Bees**. Click the key icon next to each AI you want to enable, paste in your API key from that provider's console, and save. You need at least 2 AIs for a real hive — one paid provider plus the free Gemini tier (\$5 minimum total, gets you started). For setup links and instructions per provider, see the **API Key Guide** at waxframe.com/api-details.html.

3 Pick your Builder.

The Builder is the AI that synthesizes all the reviewers' feedback every round — it does the heaviest work. On a starter hive, set the free Gemini model as Builder so you're not burning credits on the heaviest role. Builder is chosen on screen 2.

4 Set up your project.

Screens 3 through 5 collect your project goal, optional reference material (source documents the hive consults but never edits), and your starting document. The fastest path: click **Use Template** on the Project screen, pick *Start from Scratch* or *Refine an Existing Draft*, and choose a template — the goal fields populate automatically.

5 Run your first round.

Hit the round-launch button. Watch the Live Console as each reviewer's response streams in. When the round finishes, read the Conflicts panel (if any) and the new document version. Resolve conflicts by picking an option, writing your own, or keeping the original — then run another round.

Common gotchas & tips

TIP · QUICK START TEMPLATE

Brand new to WaxFrame and not sure where to begin? Pick the **Quick Start** template (a chocolate-chip cookie recipe). It runs the full hive flow on a low-stakes subject so you can see exactly how rounds, reviewers, the Builder, and convergence work end-to-end before you commit to an important document.

TIP · COST-CONSCIOUS SETUP

The minimum-viable hive (\$5 in API credits + free Gemini) works, but the jump from \$5 to ~\$10 (adding one more paid provider) gets you a 3+ AI hive that converges faster and resolves disagreements automatically. Past 3+ AIs you're in diminishing returns. Model-family diversity matters more than count — one OpenAI + one Anthropic + one Google beats three OpenAI models.

TIP · API KEYS LIVE IN YOUR BROWSER

Keys are saved to your browser's `localStorage`, per-profile per-machine. They never leave your machine except to talk to each AI provider directly (one exception: Claude routes through a tiny CloudFlare Worker because Anthropic doesn't allow direct browser calls — the Worker doesn't log keys or prompts). Clearing browser data for this site removes them. Treat them like any browser-stored credential.

TIP · STARTING FROM SCRATCH USUALLY CONVERGES FASTER

If your starting document is heavily misaligned with your goal, the hive will burn early rounds debating the existing content instead of building forward. Starting fresh from a template — even for documents you already have a draft of — often produces a better result in fewer rounds.

Where to learn more

DOCUMENT	WHAT IT COVERS
User Manual waxframe.com/waxframe-user-manual.html	Complete walk-through of every screen, the full workflow, Conflicts resolution, exporting, and the convergence model.
API Key Guide waxframe.com/api-details.html	Step-by-step API key setup for every supported provider, links to consoles, and per-provider gotchas.
Document Playbooks waxframe.com/document-playbooks.html	Proven multi-AI strategies with real convergence data for résumés, blog posts, RFP responses, business proposals, and more.
Template Catalog waxframe.com/templates.html	All 23 built-in templates: cover letters, résumés, business proposals, RFP responses, product reviews, blog posts, and others.

What Are AI Tokens?

waxframe.com/what-are-tokens.html

Plain-English explanation of token usage and what running rounds costs.

Help & Troubleshooting

waxframe.com/help.html

Diagnostic bundle export, bug reports, reset-to-known-state procedures.

Welcome to the hive

WaxFrame is open source under AGPL-3.0. Source on GitHub: github.com/WeirDave/WaxFrame-Professional. Buy a Pro license to support development: weirdave.gumroad.com/l/WaxFrame.